

साधक

Saadhak

Most AI trading agents are confident. This one is calibrated.

An autonomous, defined-risk options agent on Alpaca that earns its position size by proving it knows what it knows.

TEAM SAADHAKA • lab1aFeatherless ×
Alpaca • paper account
PA36JVGLDYZ4

The problem

Confidence is not accuracy, and nothing usually checks.

- A language model will tell you it is 90% sure.

Nothing scores that claim.

Nothing changes if it turns out to be wrong.

So the number is decoration, and the size behind it is a guess.

Saadhak makes the claim falsifiable, scores it, and lets that score decide how much money it may risk.

What it trades

Defined risk, every time.

- Iron condors and vertical spreads — they profit when the market stays inside a range.

one day to expiry, on liquid index products

never naked, never a market order

the maximum loss is known before the order goes in

Three components

Niyama

the discipline

- deterministic code

- searches 70 structures per name

- the only thing that can create an order

Saadhaka

the practitioner

- a language model on Featherless, over Alpaca's MCP server

- 33 research tools, none mutate

- may veto, cannot start a trade

Saakshya the witness — scores every probability the model states, reconciles the broker through the Alpaca CLI

Nineteen gates between a candidate and an order

Every one explains itself. Every refusal is journalled with its reason.

GROUP	GATES
market state	hours · days to expiry · reconciled · kill switch
risk appetite	whitelist · max loss · portfolio caps · loss halts · sizeable
measurement	expectancy — TP at half credit vs a 2× stop = a 2-in-3 breakeven
quality	defined risk · liquidity · delta data · idempotent id · limit only · exits
circuits	earnings guard · regime breaker · correlation

What earns the size

The mechanism the whole design rests on.

- Every review states a probability. The witness resolves it against what happened.

90% claimed and right half the time means the budget is cut

the engine sets the band, so a wider guess cannot buy a better score

Brier 0.248 over 40 · says 48%, right 57% → 0.70× of full size

It knows when to stop

- Two consecutive stop-outs block new entries in that name.
- SPY and QQQ correlate at **0.91** — one position, not two.
- A 3% daily loss halts entries; 6% flattens everything. exits keep running through a halt — it never traps a position

Every Alpaca surface, used for what it is good at

SURFACE	FOR
Trading API	multi-leg limit orders, signed negative for credit
Market Data	chains, bars, quotes, news — greeks solved locally at 0 DTE
MCP server	the practitioner's research, trading toolset withheld
Featherless	open-weight models behind an OpenAI-compatible endpoint
CLI	reconciliation, kill switch, health checks, dry-runs
Contracts	queried without an underlying, to discover what lists options

The results, honestly

- Four positions across two sessions. Three stopped out when the index trended.

the account finished at \$99,212 — down \$788

every loss capped at exactly one credit

no structure ever exceeded its stated maximum

Two sessions is not a track record. What this demonstrates is a mechanism.

What it rests on

Stated plainly, because it matters.

- Alpaca's free data tier: IEX quotes, indicative option quotes, no OPRA, 15-minute SIP delay.

every decision records its data provenance and marks degraded inputs

a stock moved 8% between our last visible print and the truth

- Calibration rests on tens of forecasts, not thousands.

Thank you

for taking the time to go through this agentic demo.

The model proposes.

The discipline decides.

A witness keeps score.

TEAM SAADHAKA

github.com/gachchami/saadhak ·

saadhak.streamlit.app

Alpaca paper account 8e02ddcf-e8fd-4b5f-
9d67-9c619b0334c9

Speaker notes